

# D3\_CrimeRate

August 9, 2018

## 1 Day-3 Use Case : Predicting Crime rate.

```
In [72]: import pandas as pd
         from sklearn import linear_model
         from sklearn.metrics import mean_squared_error, r2_score
         from sklearn.model_selection import train_test_split, cross_validate
         from math import sqrt
         import matplotlib.pyplot as plt
         import numpy as np
```

```
In [73]: df=pd.read_excel('Crime.xlsx')
```

### 1.0.1 Data Exploration

```
In [74]: df.info() # All are continuous variables
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 7 columns):
X1      50 non-null int64
X2      50 non-null int64
X3      50 non-null int64
X4      50 non-null int64
X5      50 non-null int64
X6      50 non-null int64
X7      50 non-null int64
dtypes: int64(7)
memory usage: 2.8 KB
```

```
In [75]: df.describe() # Nothing abnormal to be found
```

```
Out[75]:
```

|       | X1         | X2         | X3        | X4        | X5        | X6 \      |
|-------|------------|------------|-----------|-----------|-----------|-----------|
| count | 50.000000  | 50.000000  | 50.000000 | 50.000000 | 50.000000 | 50.000000 |
| mean  | 717.960000 | 616.180000 | 37.760000 | 58.800000 | 15.400000 | 29.900000 |
| std   | 293.938766 | 573.739175 | 13.820364 | 9.965246  | 6.023762  | 14.801062 |
| min   | 341.000000 | 29.000000  | 16.000000 | 42.000000 | 4.000000  | 7.000000  |
| 25%   | 497.000000 | 230.750000 | 30.000000 | 49.000000 | 11.000000 | 21.250000 |

|     |             |             |           |           |           |           |
|-----|-------------|-------------|-----------|-----------|-----------|-----------|
| 50% | 654.500000  | 454.000000  | 34.500000 | 59.000000 | 14.000000 | 25.000000 |
| 75% | 820.500000  | 822.500000  | 42.250000 | 67.000000 | 19.000000 | 34.250000 |
| max | 1740.000000 | 3545.000000 | 86.000000 | 81.000000 | 34.000000 | 81.000000 |

|       | X7        |
|-------|-----------|
| count | 50.000000 |
| mean  | 13.820000 |
| std   | 5.157479  |
| min   | 8.000000  |
| 25%   | 11.000000 |
| 50%   | 12.000000 |
| 75%   | 15.750000 |
| max   | 36.000000 |

```
In [76]: df.isna().sum() # No Missing Values
```

```
Out[76]: X1      0
         X2      0
         X3      0
         X4      0
         X5      0
         X6      0
         X7      0
         dtype: int64
```

```
In [77]: x=df[df.columns[1:7]]
         y=df[df.columns[0]]
```

```
In [78]: df.columns[1:7]
```

```
Out[78]: Index(['X2', 'X3', 'X4', 'X5', 'X6', 'X7'], dtype='object')
```

```
In [79]: x.sample(10)
```

```
Out[79]:
```

|    | X2   | X3 | X4 | X5 | X6 | X7 |
|----|------|----|----|----|----|----|
| 20 | 608  | 33 | 46 | 22 | 24 | 8  |
| 8  | 38   | 36 | 69 | 7  | 25 | 12 |
| 43 | 433  | 43 | 48 | 26 | 23 | 12 |
| 2  | 347  | 57 | 70 | 18 | 16 | 16 |
| 9  | 226  | 31 | 66 | 9  | 58 | 15 |
| 48 | 1022 | 82 | 72 | 22 | 15 | 16 |
| 49 | 1244 | 66 | 67 | 26 | 18 | 16 |
| 29 | 216  | 36 | 43 | 18 | 23 | 8  |
| 19 | 98   | 23 | 56 | 15 | 50 | 15 |
| 25 | 693  | 35 | 57 | 9  | 60 | 18 |

## 1.0.2 Model Preparation

```
In [97]: def testmodel(x,y):
         x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20,random_sta
```

```

regr = linear_model.LinearRegression()
regr.fit(x_train,y_train)
y_pred_test=regr.predict(x_test)
print("Root Mean squared error: %.4f"% sqrt(mean_squared_error(y_test,y_pred_test)))
print('R2 score: %.4f' % r2_score(y_test, y_pred_test))
plt.scatter(y_test,y_pred_test)


```

In [98]: testmodel(x,y)

*# Base model with all the variables.*

Root Mean squared error: 169.7250

R2 score: -0.5133



## Identifying Important Features

In [99]: df.corr()

Out [99]:


|    | X1        | X2        | X3        | X4        | X5        | X6        | X7        |
|----|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| X1 | 1.000000  | 0.756505  | 0.533198  | -0.135459 | 0.322519  | -0.175224 | -0.026283 |
| X2 | 0.756505  | 1.000000  | 0.509339  | -0.184445 | 0.291031  | -0.199070 | -0.045570 |
| X3 | 0.533198  | 0.509339  | 1.000000  | 0.120265  | 0.311526  | -0.277475 | 0.124502  |
| X4 | -0.135459 | -0.184445 | 0.120265  | 1.000000  | -0.537162 | 0.182364  | 0.681072  |
| X5 | 0.322519  | 0.291031  | 0.311526  | -0.537162 | 1.000000  | -0.626953 | -0.513958 |
| X6 | -0.175224 | -0.199070 | -0.277475 | 0.182364  | -0.626953 | 1.000000  | 0.591663  |
| X7 | -0.026283 | -0.045570 | 0.124502  | 0.681072  | -0.513958 | 0.591663  | 1.000000  |

here X2,X3,X5,X6,X4,X7 are the order of importance of variables.

```
In [100]: testmodel(x[['X2','X3','X5','X6','X4','X7']],y) # Base Model
```

Root Mean squared error: 169.7250


R2 score: -0.5133



```
In [101]: testmodel(x[['X2','X3','X5','X6','X4']],y) # Model without X7
```

Root Mean squared error: 127.5529

R2 score: 0.1453




```
In [102]: testmodel(x[['X2','X3','X5','X6']],y)
```

*# Model without X4*

Root Mean squared error: 127.7484


R2 score: 0.1427



## Identifying Outliers and removing them from Data.


```
In [103]: y_new=np.array(y)
          plt.hist(y)
```

```
Out[103]: (array([10., 12., 7., 12., 5., 0., 0., 2., 1., 1.]),
          array([ 341. , 480.9, 620.8, 760.7, 900.6, 1040.5, 1180.4, 1320.3,
                  1460.2, 1600.1, 1740. ]),
          <a list of 10 Patch objects>)
```



```
In [104]: df_rm=df[df['X1']<(np.mean(y_new)+3*np.std(y_new))]  
          x_rm=df_rm[df_rm.columns[1:7]]  
          y_rm=df_rm[df_rm.columns[0]]  
          testmodel(x_rm[['X2','X3','X5','X6']],y_rm)
```

Root Mean squared error: 256.5802  
R2 score: 0.1581



### Cross Validating the Dataset.

```
In [105]: regr=linear_model.LinearRegression()
          scores = cross_validate(regr, x_rm, y_rm,cv=10, scoring=['neg_mean_squared_error','r2']
          print(sqrt(-scores['test_neg_mean_squared_error'].mean()))
          print(scores['test_r2'].mean())
```

```
215.28213995577553
-2.431384329071694
```